

02-Visualizing-Pandas-Time-Series-Data

October 5, 2024

[]:

1 Visualizing Time Series Data

Let's go through a few key points of creating nice time series visualizations!

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[2]: df = pd.read_csv('COST.csv')
```

```
[3]: df.head()
```

```
[3]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2016-09-06	158.130005	158.149994	156.020004	158.059998	140.896622	
1	2016-09-07	157.639999	157.869995	155.399994	155.639999	138.739395	
2	2016-09-08	155.190002	155.490005	152.940002	153.470001	136.805038	
3	2016-09-09	152.589996	152.789993	150.699997	150.699997	134.335831	
4	2016-09-12	150.500000	151.990005	150.259995	151.690002	135.218338	

	Volume
0	2716900
1	2984100
2	2993100
3	2993900
4	2982400

```
[4]: df = pd.read_csv('COST.csv',index_col='Date',parse_dates=True)
```

```
[5]: df.head(10)
```

```
[5]:
```

	Open	High	Low	Close	Adj Close	\
Date						
2016-09-06	158.130005	158.149994	156.020004	158.059998	140.896622	
2016-09-07	157.639999	157.869995	155.399994	155.639999	138.739395	
2016-09-08	155.190002	155.490005	152.940002	153.470001	136.805038	

2016-09-09	152.589996	152.789993	150.699997	150.699997	134.335831
2016-09-12	150.500000	151.990005	150.259995	151.690002	135.218338
2016-09-13	151.179993	151.580002	150.070007	150.740005	134.371490
2016-09-14	150.710007	152.250000	150.199997	150.690002	134.326889
2016-09-15	150.720001	152.970001	150.110001	152.669998	136.091919
2016-09-16	151.960007	152.660004	151.139999	152.350006	135.806625
2016-09-19	152.130005	153.029999	151.600006	151.789993	135.307449


Date	Volume
2016-09-06	2716900
2016-09-07	2984100
2016-09-08	2993100
2016-09-09	2993900
2016-09-12	2982400
2016-09-13	2148600
2016-09-14	2188900
2016-09-15	2284300
2016-09-16	2842000
2016-09-19	1321300

```
[6]: # To show that dates are already parsed
df.index
```

```
[6]: DatetimeIndex(['2016-09-06', '2016-09-07', '2016-09-08', '2016-09-09',
                    '2016-09-12', '2016-09-13', '2016-09-14', '2016-09-15',
                    '2016-09-16', '2016-09-19',
                    ...,
                    '2021-08-20', '2021-08-23', '2021-08-24', '2021-08-25',
                    '2021-08-26', '2021-08-27', '2021-08-30', '2021-08-31',
                    '2021-09-01', '2021-09-02'],
                    dtype='datetime64[ns]', name='Date', length=1258, freq=None)
```


First we'll create a line plot that puts both 'Close' and 'Volume' on the same graph. Remember that we can use `df.plot()` in place of `df.plot.line()`

```
[7]: df.plot();
```




This isn't very helpful due to the difference in y-values, so we'll split them up.

```
[8]: df['Adj Close'].plot();
```



```
[9]: df['Volume'].plot();
```



2 Time Series Plot Formatting

2.1 X Limits

There are two ways we can set a specific span of time as an x-axis limit. We can plot a slice of the dataset, or we can pass x-limit values as an argument into `df.plot()`.


The advantage of using a slice is that pandas automatically adjusts the y-limits accordingly.

The advantage of passing in arguments is that pandas automatically tightens the x-axis. Plus, if we're also setting y-limits this can improve readability.

2.1.1 Choosing X Limits by Slice:


```
[10]: # Dates are separated by a colon:
      df['Close']['2018-01-01':'2019-01-01'].plot(figsize=(12,4))
```

```
[10]: <Axes: xlabel='Date'>
```



2.1.2 Choosing X Limits by Argument:

```
[11]: # Why is the y-axis so different?
# Because pandas first plotted the whole thing, then it narrowed
# down the plot, unlike above, where we narrowed the df first!
df['Close'].plot(figsize=(12,4),xlim=['2018-01-01','2019-01-01']);
```



NOTE: It's worth noting that the limit values do not have to appear in the index. Pandas will plot the actual dates based on their location in time. Also, another advantage of slicing over arguments is that it's easier to include the upper/lower bound as a limit. That is, `df['column']['2017-01-01:'].plot()` is easier to type than `df['column'].plot(xlim=('2017-01-01',df.index.max()))`

Now let's focus on the y-axis limits to get a better sense of the shape of the data. First we'll find out what upper and lower limits to use.


```
[12]: # FIND THE MINIMUM VALUE IN THE RANGE:
df.loc['2018-01-01':'2019-01-01']['Close'].min()
```

```
[12]: 178.610001
```

```
[13]: # FIND THE MAXIMUM VALUE IN THE RANGE:  
df.loc['2018-01-01':'2019-01-01']['Close'].max()
```


```
[13]: 244.210007
```

```
[14]: # PLUG THESE IN AS Y-LIMIT VALUES:  
df['Close'].  
    plot(figsize=(12,4),xlim=['2018-01-01','2019-01-01'],ylim=[170,250]);
```



NOTE: Be careful when setting y-axis limits! Setting too narrow a slice can make graphs appear overly volatile.

```
[15]: df['Close'].plot(figsize=(12,4),xlim=['2018-01-01','2019-01-01'],ylim=[0,600]);
```



2.2 X Ticks

In this section we'll look at how to change the format and appearance of dates along the x-axis. To do this, we'll borrow a tool from matplotlib called `dates`.


```
[16]: from matplotlib import dates
```

2.2.1 Set the spacing


The x-axis values can be divided into major and minor axes. For now, we'll work only with the major axis and learn how to set the spacing with `.set_major_locator()`.

```
[17]: #Notice the strange "jump" in the middle due to the month change!  
df['Close'].plot(xlim=['2018-01-01','2018-03-01'],ylim=(170,210))
```

```
[17]: <Axes: xlabel='Date'>
```



```
[18]: # you can remove "date" index name label  
ax = df['Close'].plot(xlim=['2018-01-01','2018-03-01'],ylim=(170,210),  
                      xlabel='',ylabel='price')
```





2.2.2 Using Tick Locator


```
[19]: from matplotlib import dates
```

```
[20]: ax = df['Close'].plot(xlim=['2018-01-01','2018-03-01'],ylim=(170,210),  
                           xlabel='',ylabel='price')
```

```
# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS  
ax.xaxis.set_major_locator(dates.WeekdayLocator())
```




```
[21]: ax = df['Close'].plot(xlim=['2018-01-01', '2018-03-01'], ylim=(170, 210),  
                           xlabel='', ylabel='price')  
  
# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS  
ax.xaxis.set_major_locator(dates.WeekdayLocator(byweekday=0))
```



Notice that dates are spaced one week apart. The dates themselves correspond with `byweek-day=0`, or Mondays. For a full list of locator options available from `matplotlib.dates` visit https://matplotlib.org/api/dates_api.html#date-tickers


```
[22]: ax = df['Close'].plot(xlim=['2018-01-01', '2018-03-01'], ylim=(170, 210),
                           xlabel='', ylabel='price')

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.MonthLocator())
```



```
[23]: ax = df['Close'].plot(xlim=['2018-01-01','2018-03-01'],ylim=(170,210),
                           xlabel='',ylabel='price')

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.YearLocator()) # DayLocator()
```



2.3 Formator

2.3.1 Date Formatting

Formatting follows the Python datetime strftime codes. The following examples are based on `datetime.datetime(2001, 2, 3, 16, 5, 6)`:

CODE

MEANING

EXAMPLE

`%Y`

Year with century as a decimal number.

2001

`%y`

Year without century as a zero-padded decimal number.

01

`%m`

Month as a zero-padded decimal number.

02

%B

Month as locale's full name.

February

%b

Month as locale's abbreviated name.

Feb

%d

Day of the month as a zero-padded decimal number.

03

%A

Weekday as locale's full name.

Saturday

%a

Weekday as locale's abbreviated name.

Sat

%H

Hour (24-hour clock) as a zero-padded decimal number.

16

%I

Hour (12-hour clock) as a zero-padded decimal number.

04

%p

Locale's equivalent of either AM or PM.

PM

%M

Minute as a zero-padded decimal number.

05

%S

Second as a zero-padded decimal number.

06

CODE

MEANING

EXAMPLE

`%#m`

Month as a decimal number. (Windows)

2

`%-m`

Month as a decimal number. (Mac/Linux)

2

`%#x`

Long date

Saturday, February 03, 2001

`%#c`

Long date and time

Saturday, February 03, 2001 16:05:06

```
[24]: # USE THIS SPACE TO EXPERIMENT WITH DIFFERENT FORMATS
from datetime import datetime
datetime(2001, 2, 3, 16, 5, 6).strftime("%A, %B %d, %Y %I:%M:%S %p")
```


```
[24]: 'Saturday, February 03, 2001 04:05:06 PM'
```

2.3.2 Combine Locator with Formatter

```
[25]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close'].plot(xlim=['2018-01-01','2018-03-01'],ylim=(170,210),
                    xlabel='',ylabel='price')

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.WeekdayLocator(byweekday=0))

# ADD IN THE FORMATTER
ax.xaxis.set_major_formatter(dates.DateFormatter("%a-%B-%d"))
```




2.4 Major vs. Minor Axis Values

All of the tick marks we've used so far have belonged to the major axis. We can assign another level called the minor axis, perhaps to separate month names from days of the month.

DEFAULT PLOT


```
[26]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close']['2018-01-01':'2019-03-01'].plot()
```



MAJOR TICK PLOT EDITS

```
[27]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close']['2018-01-01':'2020-01-01'].plot()

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.YearLocator())
ax.xaxis.set_major_formatter(dates.DateFormatter("%Y"))
```





MAJOR AND MINOR TICK PLOT EDITS

```
[28]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close']['2018-01-01':'2020-01-01'].plot()

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.YearLocator())
ax.xaxis.set_major_formatter(dates.DateFormatter("%Y"))

# SET THE TICK LOCATOR AND FORMATTER FOR THE MINOR AXIS
ax.xaxis.set_minor_locator(dates.MonthLocator())
ax.xaxis.set_minor_formatter(dates.DateFormatter("%B"))
```



Rotate Minor Ticks


```
[29]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close']['2018-01-01':'2020-01-01'].plot()

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.YearLocator())
ax.xaxis.set_major_formatter(dates.DateFormatter("%Y"))
```

```
# SET THE TICK LOCATOR AND FORMATTER FOR THE MINOR AXIS
ax.xaxis.set_minor_locator(dates.MonthLocator())
ax.xaxis.set_minor_formatter(dates.DateFormatter("%B"))

# FURTHER EDITING TICK PROPERTIES

ax.tick_params(axis="x", which="major", rotation=90,pad=50)
ax.tick_params(axis="x", which="minor", rotation=90)
```



Include January instead of just Year on Major Tick

```
[30]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close']['2018-01-01':'2020-01-01'].plot()

# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.YearLocator())
ax.xaxis.set_major_formatter(dates.DateFormatter("%Y %B"))


# SET THE TICK LOCATOR AND FORMATTER FOR THE MINOR AXIS
ax.xaxis.set_minor_locator(dates.MonthLocator())
ax.xaxis.set_minor_formatter(dates.DateFormatter("%B"))

# FURTHER EDITING TICK PROPERTIES

ax.tick_params(axis="x", which="major",rotation=90, pad=5)
ax.tick_params(axis="x", which="minor",rotation=90)

# To get it exactly perfect is a lot more work:
```

```
# https://stackoverflow.com/questions/28615887/
# how-to-move-a-ticks-label-in-matplotlib
plt.xticks(ha='center');
```



2.5 Adding Gridlines

We can add x and y axis gridlines that extend into the plot from each major tick mark.

```
[31]: plt.figure(dpi=100,figsize=(10,3))
ax = df['Close']['2018-01-01':'2020-01-01'].plot()


# SET THE TICK LOCATOR AND FORMATTER FOR THE MAJOR AXIS
ax.xaxis.set_major_locator(dates.YearLocator())
ax.xaxis.set_major_formatter(dates.DateFormatter("%Y %B"))

# SET THE TICK LOCATOR AND FORMATTER FOR THE MINOR AXIS
ax.xaxis.set_minor_locator(dates.MonthLocator())
ax.xaxis.set_minor_formatter(dates.DateFormatter("%B"))

# FURTHER EDITING TICK PROPERTIES

ax.tick_params(axis="x", which="major",rotation=90, pad=5)
ax.tick_params(axis="x", which="minor",rotation=90)

ax.yaxis.grid(True)
ax.xaxis.grid(True)
```



[]: